

<b>Name of the student:</b>	Tanmay Prashant Rane	<b>Roll No.</b>	8031
<b>Practical Number:</b>	4	<b>Date of Practical:</b>	
<b>Relevant CO's</b>	<b>At the end of the course students will be able to use tools like hadoop and NoSQL to solve big data related problems.</b>		
<b>Sign here to indicate that you have read all the relevant material provided before attempting this practical</b>			<b>Sign:</b>

### Practical grading using Rubrics

Indicator	Very Poor	Poor	Average	Good	Excellent
<b>Timeline</b> (2)	More than a session late (0)	NA	NA	NA	Early or on time (2)
<b>Code de- sign</b> (2)	N/A	Very poor code design with no comments and indentation(0.5)	Poor code design with very comments and indentation (1)	Design with good coding standards (1.5)	Accurate design with better coding standards (2)
<b>Performance</b> (4)	Unable to perform the experiment (0)	Able to partially perform the experiment (1)	Able to perform the experiment for certain use cases (2)	Able to perform the experiment considering most of the use cases (3)	Able to perform the experiment considering all use cases (4)
<b>Postlab</b> (2)	No Execution(0)	N/A	Partially Executed (1)	N/A	Fully Executed (2)


Total Marks (10)	Sign of instructor with date

# Practical

Course title: Big Data Analytics  
 Course term: 2019-2020  
 Instructor name: Saurabh Kulkarni

**Problem Statement: Perform matrix multiplication using one step map-reduce**

**Theory: Explain the concept of matrix multiplication using one step map-reduce with the help of an example**

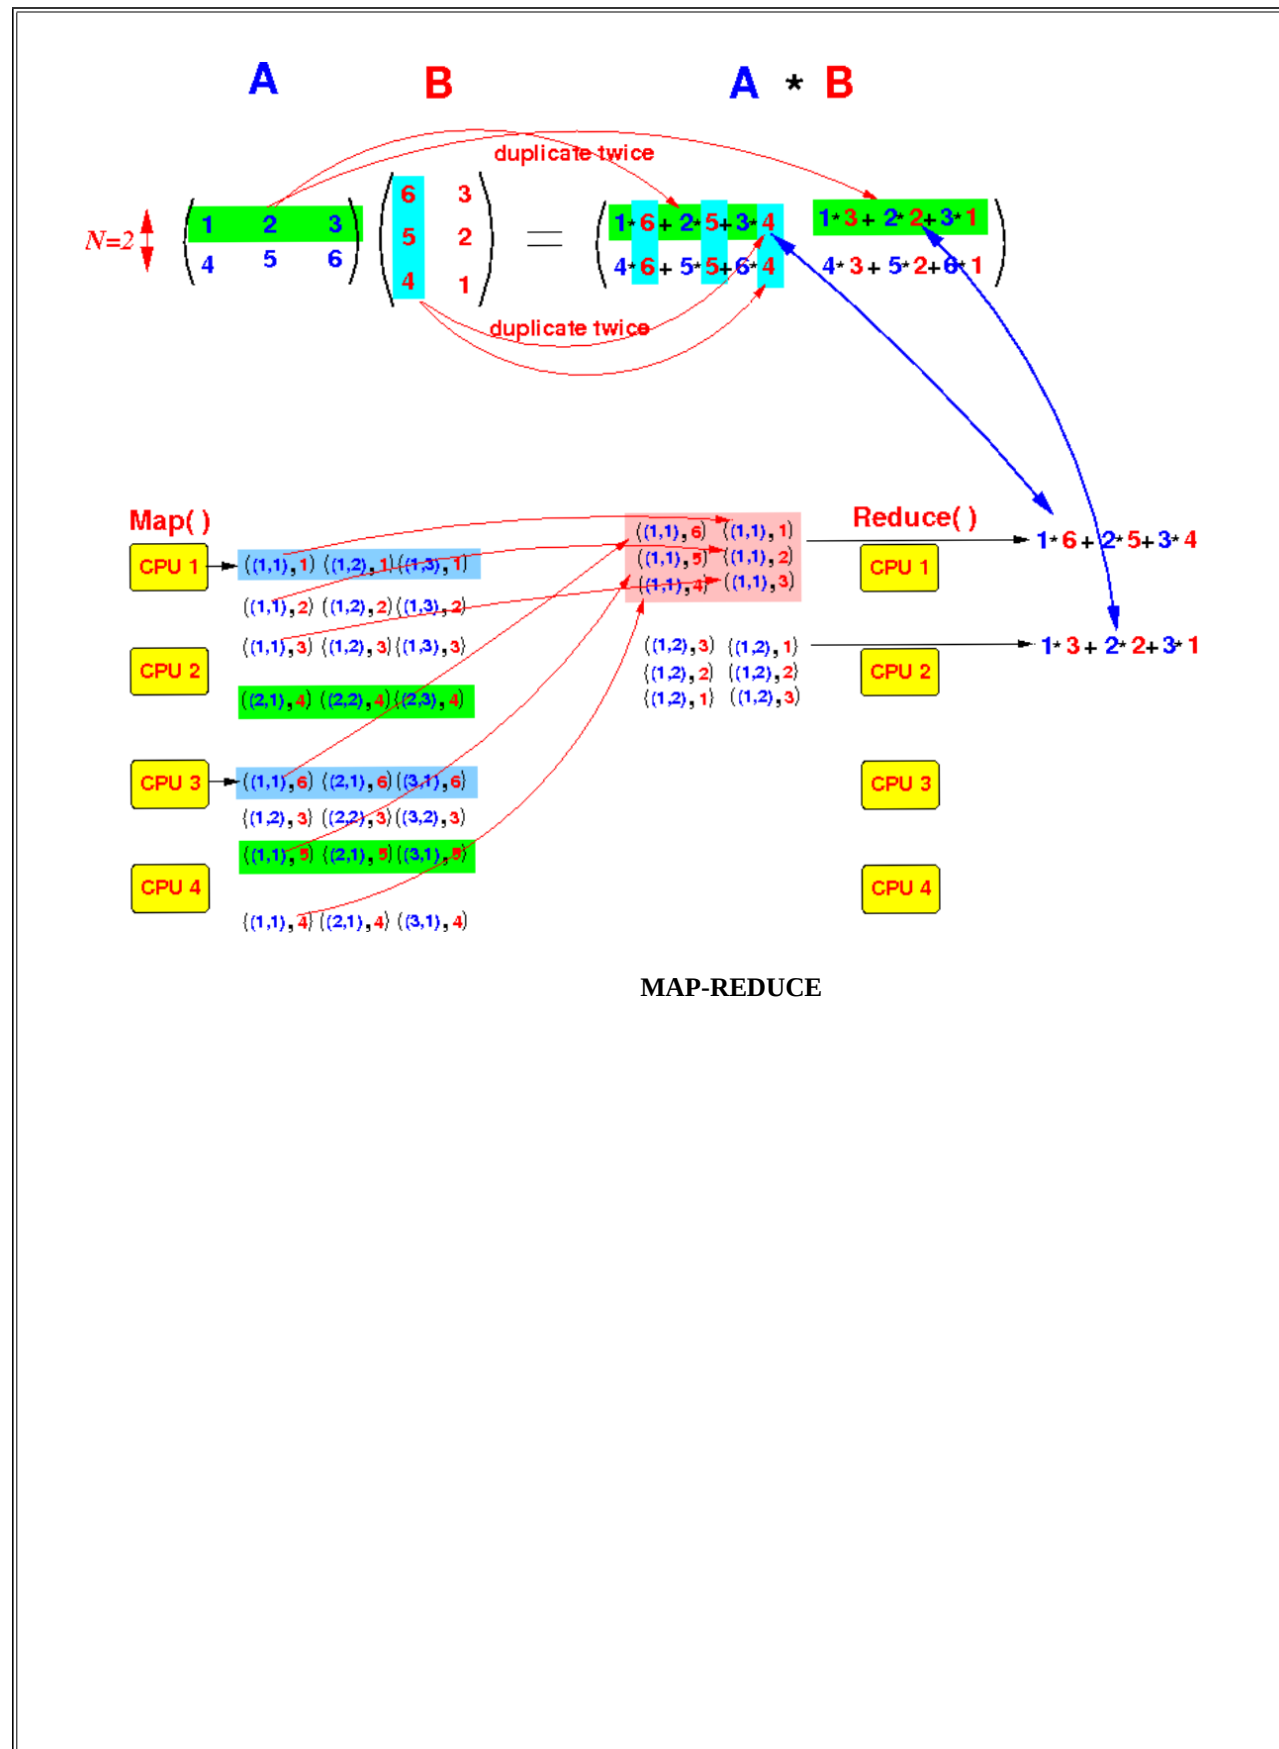


MAP

$$\begin{matrix} \text{A} & & \text{B} & & & & \text{A} * \text{B} \\ \begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{pmatrix} & & \begin{pmatrix} 6 & 3 \\ 5 & 2 \\ 4 & 1 \end{pmatrix} & = & \begin{pmatrix} 1*6 + 2*5 + 3*4 & 1*3 + 2*2 + 3*1 \\ 4*6 + 5*5 + 6*4 & 4*3 + 5*2 + 6*1 \end{pmatrix} \end{matrix}$$

Input to "Matrix Multiplication"

INPUT



**Code:****code for mapper:****Code for Reducer:****Code for Driver Class:**

```

import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;

public class Map extends Mapper<LongWritable, Text, Text, Text>
{
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException
    {
        //Create a configuration object using context
        Configuration conf = context.getConfiguration();

        // Consider a matrix of size m*n and other matrix is of size n*p. Get the values of m and p using
        get() method of configuration object
        int m = Integer.parseInt(conf.get("m"));
        int p = Integer.parseInt(conf.get("p"));

        //convert Text value to string
        String line = value.toString();

        // Split each line into tokens separated by comma.
        String[] indicesAndValue = line.split(",");

        // Define outputkey of type Text()
        Text outputKey = new Text();

        //Define outputvalue of type Text()
        Text outputValue = new Text();

        //If first token is A
        if (indicesAndValue[0].equals("A"))
        {
            // Vary k from 0 to no. of columns-1 of 2nd matrix
            for (int k=0; k < p; k++) {
                //set outputkey as tokens[1],k i.e.i,k
                outputKey.set(indicesAndValue[1] + "," + k);
                // set outputkeyvalue as A,tokens[2],tokens[3] i.e. A,j,mij
            }
        }
    }
}

```

```

outputValue.set(indicesAndValue[0] + "," + indicesAndValue[2] + "," + indicesAndValue[3]);
//outputValue
context.write(outputKey, outputValue);

}
}

// Vary i from 0 to no.of rows-1 of 1st matrix
for (int i=0; i < m; i++)
{
//set outputkey as i,toknes[2] i.e.i,k
outputKey.set(i + "," + indicesAndValue[2]);

//set outputkey value B,toknes[1],tokens[3] i.e. B,j,nkj
outputValue.set("N," + indicesAndValue[1] + "," + indicesAndValue[3]);
// write key and value to context
context.write(outputKey, outputValue);
}
}
}

```

#### //REDUCER

```

import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;

public class Reduce extends Reducer<Text, Text, Text, Text> {

    public void reduce(Text key, Iterable<Text> values, Context context)
        throws IOException, InterruptedException {
        // process values

        // define value as String type array
        String[] value;
        //create 2 hashmaps (hashA,hasB) of Integer,Float to store values from A and B matrix
        HashMap<Integer, Float> hashA = new HashMap<Integer, Float>();
        HashMap<Integer, Float> hashB = new HashMap<Integer, Float>();
        //for each value in values
        for (Text val : values) {
            // convert each value to string as it is Text and split it by comma. Store this in value i.e. string type array
            // defined above.
            value = val.toString().split(",");
            //if value[0] is A then
            if(value[0].equals("A")) {
                hashA.put(Integer.parseInt(value[1]), Float.parseFloat(value[2]));
            }

            //put value[1] i.e.j and value[2] i.e. mij in hashA

```

```

else
    {
        hashB.put(Integer.parseInt(value[1]), Float.parseFloat(value[2]));
    }
}

//else put value[1] i.e.j and value[2] i.e. njk in hashB

// take value of n which is common in both the matrices here from context object
int n = Integer.parseInt(context.getConfiguration().get("n"));
// define result of type float and assign value 0 to it.
float result = 0.0f;
float m_ij;// define a_ij of type float
float n_jk;// define b_jk of type float

//define a loop variable j and iterate till it is less than n
for (int j = 0; j < n; j++)
{
    // check if value exists for a key j in hashA. if yes assign it to a_ij else assign 0 to a_ij.
    m_ij = hashA.containsKey(j) ? hashA.get(j) : 0.0f;
    // check if value exists for a key j in hashB. if yes assign it to b_jk else assign 0 to b_jk.
    n_jk = hashB.containsKey(j) ? hashB.get(j) : 0.0f;
    // multiply a_ij with b_jk and add this product to result which is of type float and declared above.
    result += m_ij * n_jk;

}
// if value of result is not 0 then
if (result != 0.0f)
{
    // write key,value to context. key is null and value is (i,k,result)
    context.write(null, new Text(key.toString() + "," + Float.toString(result)));
}
}
}

```

#### //DRIVER

```

import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
public class MatDriver {

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        // A is an m-by-n matrix; B is an n-by-p matrix.
        conf.set("m", "3");
        conf.set("n", "3");
        conf.set("p", "3");
        @SuppressWarnings("deprecation")
        Job job = new Job(conf, "MatrixMatrixMultiplicationOneStep");
        job.setJarByClass(MatDriver.class);
        job.setOutputKeyClass(Text.class);
    }
}

```

```
job.setOutputValueClass(Text.class);

    job.setMapperClass(Map.class);
    job.setReducerClass(Reduce.class);

    job.setInputFormatClass(TextInputFormat.class);
    job.setOutputFormatClass(TextOutputFormat.class);

    FileInputFormat.addInputPath(job, new Path("/media/tanmay/Data/SEM-8/BDA/EXP4/2by2data"));
    FileOutputFormat.setOutputPath(job, new Path("/media/tanmay/Data/SEM-8/BDA/EXP4/MatOut3"));
    job.waitForCompletion(true);
}
}
```

**PostLab:**

1. Generate a 100x100 matrix in the format that above map-reduce code understands using su programming language
2. Compute execution time for a 100x100 matrix using suitable APIs

**Code for postlab question****//DRIVER FOR 100 X 100 MATRIX MULTIPLICATION**

```

import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
public class MatDriver {

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        // A is an m-by-n matrix; B is an n-by-p matrix.
        long startTime = System.nanoTime();
        conf.set("m", "100");
        conf.set("n", "100");
        conf.set("p", "100");
        @SuppressWarnings("deprecation")
        Job job = new Job(conf, "MatrixMatrixMultiplicationOneStep");
        job.setJarByClass(MatDriver.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(Text.class);

        job.setMapperClass(Map.class);
        job.setReducerClass(Reduce.class);

        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);

        FileInputFormat.addInputPath(job, new Path("/media/tanmay/Data/SEM-8/BDA/EXP4/post_input_1"));
        FileOutputFormat.setOutputPath(job, new Path("/media/tanmay/Data/SEM-8/BDA/EXP4/postout"));
        if (job.waitForCompletion(true)) {
            long endTime = System.nanoTime();
            long timeElapsed = endTime - startTime;
            System.out.println("Execution time in milliseconds : " + timeElapsed/1000000);
            System.out.println("Job Completed");
            return;
        }
    }
}

//EXECUTION TIME : 15136 ms ~ 15.136 s

```