

<b>Name of the student:</b>	Tanmay Prashant Rane	<b>Roll No.</b>	8031
<b>Practical Number:</b>	10	<b>Date of Practical:</b>	
<b>Relevant CO's</b>	At the end of the course students will be able to apply big data analytics in real life applications.		
<b>Sign here to indicate that you have read all the relevant material provided before attempting this practical</b>			<b>Sign:</b>

Practical grading using Rubrics Practical grading using Rubrics

Indicator	Very Poor	Poor	Average	Good	Excellent
<b>Timeline</b> (2)	More than a session late (0)	NA	NA	NA	Early or on time (2)
<b>Code de-sign</b> (2)	N/A	Very poor code design with no comments and indentation(0.5)	Poor code design with very comments and indentation (1)	Design with good coding standards (1.5)	Accurate design with better coding standards (2)
<b>Performance</b> (4)	Unable to perform the experiment (0)	Able to partially perform the experiment (1)	Able to perform the experiment for certain use cases (2)	Able to perform the experiment considering most of the use cases (3)	Able to perform the experiment considering all use cases (4)
<b>Postlab</b> (2)	Both answers incorrect (0)	N/A	One answer correct (1)	N/A	Both answers correct (2)

<b>Total Marks (10)</b>	<b>Sign of instructor with date</b>

# Practical

COURSE TITLE: BIG DATA ANALYTICS

COURSE TERM: 2019-2020

---

**Problem Statement: To find common friends in social network graph using map-reduce.**

---

## Theory:

---

Given a social network with tens of millions of users, in this chapter we'll implement a MapReduce program to identify "common friends" among all pairs of users. Let  $U$  be a set of all users:  $U_1, U_2, \dots, U_n$ . Our goal is to find common friends for every pair of  $(U_i, U_j)$  where  $i$  not equal to  $j$ .

These days most social network sites (such as Facebook, hi5, and LinkedIn) offer services to help you share messages, pictures, and videos among friends. Some sites even offer video chat services to help you connect with friends. By definition, a friend is a person whom one knows, likes, and trusts. For example, Facebook has your list of friends, and friend relationships are bidirectional on the site; if I'm your friend, you're mine too. Typically social networks precompute calculations when they can to reduce the processing time of requests, and one common processing request involves the "You and Mary (your friend) have 185 friends in common" feature. When you visit someone's profile, you see a list of friends that you have in common. This list doesn't change frequently, so it is wasteful for the site to recalculate it every time you visit that person's profile.

### Input

We prepare input as a set of records, where each record has the following format:

`<person><,><friend 1 ><friend 2 >...<friend N >`

where `<friend 1 ><friend 2 > ... <friend N >` are friends of the `<person>`.

Note that in real projects/applications, each person/friend will be identified as a unique user ID. A very simple and complete example of input is as follows:

```
100, 200 300 400 500 600
200, 100 300 400
300, 100 200 400 500
400, 100 200 300
500, 100 300
600, 100
```

In this example, user 500 has two friends identified by the user IDs 100 and 300, and user 600 has only one friend: user 100.

The mapper accepts a (key 1, value 1) pair, where key 1 is a person and value 1 is a list of associated friends of that person. The mapper emits a set of new (key 2, value 2) pairs; key 2 is a Tuple2(key 1, friend i), where friend i belongs value 1, and value 2 is the same as value 1 (a list of all friends for key 1). The reducer's key is a pair of two users (User j, User k) and value is a list of sets of friends. The reduce() function will intersect all sets of friends to find common and mutual friends for the (User j, User k) pair.

```
1 // key is the person
2 // value is a list of friends for this key=person
3 // value = (<friend_1> <friend_2> ... <friend_N>)
4 map(key, value) {
5     reducerValue = (<friend_1>< friend_2> ...< friend_N>);
6     foreach friend in (<friend_1>< friend_2> ... <friend_N>) {
7         reducerKey = buildSortedKey(person, friend);
```

```

8     emit(reducerKey, reducerValue);
9 }
10 }

```

```

1 Tuple2 buildSortedKey(person1, person2) {
2     if (person1 < person2) {
3         return new Tuple2(person1, person2);
4     }
5     else {
6         return new Tuple2(person2, person1);
7     }
8 }

```

```

1 // key = Tuple2(person1, person2)
2 // value = List {List_1, List_2, ..., List_M}
3 //   where each List_i = { set of unique user IDs }
4 reduce(key, value) {
5     outputKey = key;
6     outputValue = intersection (List_1, List_1, ..., List_M);
7     emit (outputKey, outputValue);
8 }

```

Mapreduce in action for given example

map(100, (200 300 400 500 600)) will generate:

```

([100, 200], [200 300 400 500 600])
([100, 300], [200 300 400 500 600])
([100, 400], [200 300 400 500 600])
([100, 500], [200 300 400 500 600])
([100, 600], [200 300 400 500 600])

```

map(200, (100 300 400)) will generate:

```
([100, 200], [100 300 400])  
([200, 300], [100 300 400])  
([200, 400], [100 300 400])
```

map(300, (100 200 400 500)) will generate:

```
([100, 300], [100 200 400 500])  
([200, 300], [100 200 400 500])  
([300, 400], [100 200 400 500])  
([300, 500], [100 200 400 500])
```

map(400, (100 200 300)) will generate:

```
([100, 400], [100 200 300])  
([200, 400], [100 200 300])  
([300, 400], [100 200 300])
```

map(500, (100 300)) will generate:

```
([100, 500], [100 300])  
([300, 500], [100 300])
```

map(600, (100)) will generate:

```
([100, 600], [100])
```

So, the mappers generate the following key-value pairs:

```
([100, 200], [200 300 400 500 600])
([100, 300], [200 300 400 500 600])
([100, 400], [200 300 400 500 600])
([100, 500], [200 300 400 500 600])
([100, 600], [200 300 400 500 600])
([100, 200], [100 300 400])
([200, 300], [100 300 400])
([200, 400], [100 300 400])
([100, 300], [100 200 400 500])
([200, 300], [100 200 400 500])
([300, 400], [100 200 400 500])
([300, 500], [100 200 400 500])
([100, 400], [100 200 300])
([200, 400], [100 200 300])
([300, 400], [100 200 300])
([100, 500], [100 300])
([300, 500], [100 300])
([100, 600], [100])
```

Before these key-value pairs are sent to the reducers, they are grouped by keys:

```
([100, 200], [200 300 400 500 600])
([100, 200], [100 300 400])
=> ([100, 200], ([200 300 400 500 600], [100 300 400]))

([100, 300], [200 300 400 500 600])
([100, 300], [100 200 400 500])
=> ([100, 300], ([200 300 400 500 600], [100 200 400 500]))
```

So, the reducers will receive the following set of key-value pairs:

```
([100, 200], ([200 300 400 500 600], [100 300 400]))  
([100, 300], ([200 300 400 500 600],[100 200 400 500]))  
([100, 400], ([200 300 400 500 600], [100 200 300]))  
([100, 500], ([200 300 400 500 600], [100 300]))  
([200, 300], ([100 300 400],[100 200 400 500]))  
([200, 400], ([100 300 400],[100 200 300]))  
([300, 400], ([100 200 400 500][100 200 300]))  
([300, 500], ([100 200 400 500],[100 300]))  
([100, 600], ([200 300 400 500 600], [100]))
```

Finally, the reducers will generate:

```
([100, 200], [300, 400])  
([100, 300], [200, 400, 500])  
([100, 400], [200, 300])  
([100, 500], [300])  
([200, 300], [100, 400])  
([200, 400], [100, 300])  
([300, 400], [100, 200])  
([300, 500], [100])  
([100, 600], [])
```

Following the generated output, we can see that users 100 and 600 have no common friends. The business case is that when user 100 visits user 200's profile, we can now quickly look up the [100, 200] key and see that they have two friends in common: [300, 400] . Meanwhile, the users identified by 100 and 500 have one friend (identified by user ID 300) in common.

**Code:**

Write map-reduce code to implement algorithm

**code for mapper:**

```
import java.io.IOException;
import java.util.Arrays;
import java.util.HashMap;
import java.util.List;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;

public class Map extends Mapper<LongWritable, Text, Text, Text> {

    private Text word = new Text();

    public void map(LongWritable key, Text value, Context context) throws
        ↪ IOException, InterruptedException {
        String[] line = value.toString().split(",_");
        if (line.length == 2) {
            String friend1 = line[0];
            List<String> values = Arrays.asList(line[1].split("_"));
            for (String friend2 : values) {
                int f1 = Integer.parseInt(friend1);
                int f2 = Integer.parseInt(friend2);
                if (f1 < f2)
                    word.set(friend1 + "," + friend2);
                else
                    word.set(friend2 + "," + friend1);
                context.write(word, new Text(line[1]));
            }
        }
    }
}
```



**Code for Reducer:**

```

public class Reduce extends Reducer<Text, Text, Text, Text> {
private Text result = new Text();

public void reduce(Text key, Iterable<Text> values, Context context)
    ↪ throws IOException, InterruptedException {
    HashMap<String, Integer> map = new HashMap<String, Integer>();
    StringBuilder sb = new StringBuilder();
    for (Text friends : values) {
        List<String> temp = Arrays.asList(friends.toString().split("_"));
        for (String friend : temp) {
            if (map.containsKey(friend))
                sb.append(friend + ',');
            else
                map.put(friend, 1);
        }
    }
    if (sb.lastIndexOf(",") > -1) {
        sb.deleteCharAt(sb.lastIndexOf(","));
    }

    result.set(new Text(sb.toString()));
    context.write(key, result);
}

```

**Code for Driver Class:**

```

public class MutualFriends {

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    String[] otherArgs = new GenericOptionsParser(conf, args).
        ↪ getRemainingArgs();
    // get all args
    //          if (otherArgs.length != 2) {
    //              System.err.println("Usage: Mutual Friend <
    ↪ inputfile hdfs path> <output file hdfs path>");
    //              System.exit(2);
    //          }

    @SuppressWarnings("deprecation")
    Job job = new Job(conf, "MutualFriend");
    job.setJarByClass(MutualFriends.class);
    job.setMapperClass(Map.class);
    job.setReducerClass(Reduce.class);
}

```

```
job.setOutputKeyClass(Text.class);

job.setOutputValueClass(Text.class);
// set the HDFS path of the input data
FileInputFormat.addInputPath(job, new Path("/media/tanmay/Data/SEM-8/
    ↪ BDA/EXP10/bdaexp10/BDAexp10data"));
// set the HDFS path for the output
FileOutputFormat.setOutputPath(job, new Path("/media/tanmay/Data/SEM-8/
    ↪ BDA/EXP10/bdaexp10/out"));
// Wait till job completion
System.exit(job.waitForCompletion(true) ? 0 : 1);
}

}
```

**Output:**

100,200	300,400
100,300	200,400,500
100,400	200,300
100,500	300
100,600	
200,300	100,400
200,400	100,300
300,400	100,200
300,500	100

**PostLab:**

Explain applications of above code in Facebook as social network

**Answer for postlab question**

Facebook is the latest in a long line of what we now know as “social networking” websites. But what sets it apart from the competitors, is its popularity. At last check, Facebook boasts over 2.23 billion active users.

With a billion users and requirements to analyze more than 105 terabytes every 30 minutes, Facebook’s appetite for crunching data has reached very huge proportions.

Hadoop is the key tool Facebook uses, not simply for analysis, but as an engine to power many features of the Facebook site, including messaging. That multitude of monster workloads drove the company to launch its Prism project, which supports geographically distributed Hadoop data stores.

Applications of above code for Facebook are as follows:

1. Finding mutual friends of two users instantly
2. Establishing suggestions for friendship on basis of mutual friends
3. Suggesting adds that mutual friends see to the group of friends
4. Constructing the "you may know section"
5. Pushing activities to feed that happened in cluster of particular group
6. Suggesting games and online challenges

Explain applications of above code in LinkedIn as social network

**Answer for postlab question**

With more than 400 million profiles (122 million in US and 33 million in India) across 200+ countries, more than 100 million unique monthly visitors, 3 million company pages, 2 new members joining the network every second, 5.7 billion professional searches in 2012, 7600 full-time employees, \$780 million revenue as of Oct, 2015 and earnings of 78 cents per share LinkedIn is the largest social network for professionals. People prefer to share their expertise and connect with like-minded professionals to discuss various issues of interest in a platform like LinkedIn, as it allows them to represent themselves formally in a less traditional manner. 2 or more people join LinkedIn's professional network every second, making up the pool of 400 million members

1. People you may know
2. Jobs you maybe interested in
3. News Feed Updates
4. Skill Endorsements
5. Pushing activities to feed that happened in cluster of particular group
6. Suggesting games and online challenges